

Model Building

Team ID	PNT2022TMID12635
Project Name	Machine Learning based Vehicle Performance Analyzer

Importing Libraries

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula.api as smf
```

Importing Dataset

In [2]:

```
dataset=pd.read_csv('/content/car performance.csv')
dataset
```

Out[2]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18.0	8	307.0	130	3504	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	ford torino
...
393	27.0	4	140.0	86	2790	15.6	82	1	ford mustang gl
394	44.0	4	97.0	52	2130	24.6	82	2	vw pickup
395	32.0	4	135.0	84	2295	11.6	82	1	dodge rampage
396	28.0	4	120.0	79	2625	18.6	82	1	ford ranger
397	31.0	4	119.0	82	2720	19.4	82	1	chevy s-10

398 rows x 9 columns

Finding missing data

In [3]:

```
dataset.isnull().any()
```

Out[3]:

```
mpg                False
cylinders           False
displacement        False
horsepower          False
weight              False
acceleration        False
model year          False
origin              False
car name            False
dtype: bool
```

There are no null characters in the columns but there is a special character '?' in the 'horsepower' column. So we we replaced '?' with nan and replaced nan values with mean of the column.

In [4]:

```
dataset['horsepower']=dataset['horsepower'].replace('?',np.nan)
```

In [5]:

```
dataset['horsepower'].isnull().sum()
```

Out[5]:

6

In [6]:

```
dataset['horsepower']=dataset['horsepower'].astype('float64')
```

In [7]:

```
dataset['horsepower'].fillna((dataset['horsepower'].mean()),inplace=True)
```

In [8]:

```
dataset.isnull().any()
```

Out[8]:

```
mpg                False
cylinders           False
displacement        False
horsepower           False
weight              False
acceleration         False
model year          False
origin              False
car name            False
dtype: bool
```

In [9]:

```
dataset.info() #Pandas dataframe.info() function is used to get a quick overview of the dataset.
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   mpg             398 non-null   float64
 1   cylinders        398 non-null   int64
 2   displacement     398 non-null   float64
 3   horsepower       398 non-null   float64
 4   weight           398 non-null   int64
 5   acceleration     398 non-null   float64
 6   model year       398 non-null   int64
 7   origin           398 non-null   int64
 8   car name         398 non-null   object
dtypes: float64(4), int64(4), object(1)
memory usage: 28.1+ KB
```

In [10]:

```
dataset.describe() #Pandas describe() is used to view some basic statistical details of a data frame or a series of numeric values.
```

Out[10]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin
count	398.000000	398.000000	398.000000	398.000000	398.000000	398.000000	398.000000	398.000000
mean	23.514573	5.454774	193.425879	104.469388	2970.424623	15.568090	76.010050	1.572864

std	7.815984 mpg	1.701004 cylinders	104.269838 displacement	38.199187 horsepower	846.841774 weight	2.757689 acceleration	3.697627 model year	0.802055 origin
min	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000	70.000000	1.000000
25%	17.500000	4.000000	104.250000	76.000000	2223.750000	13.825000	73.000000	1.000000
50%	23.000000	4.000000	148.500000	95.000000	2803.500000	15.500000	76.000000	1.000000
75%	29.000000	8.000000	262.000000	125.000000	3608.000000	17.175000	79.000000	2.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000	24.800000	82.000000	3.000000

There is no use with car name attribute so drop it

In [11]:

```
dataset=dataset.drop('car name',axis=1) #dropping the unwanted column.
```

In [12]:

```
corr_table=dataset.corr()#Pandas dataframe.corr() is used to find the pairwise correlatio  
n of all columns in the dataframe.  
corr_table
```

Out[12]:


	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin
mpg	1.000000	-0.775396	-0.804203	-0.771437	-0.831741	0.420289	0.579267	0.563450
cylinders	-0.775396	1.000000	0.950721	0.838939	0.896017	-0.505419	-0.348746	-0.562543
displacement	-0.804203	0.950721	1.000000	0.893646	0.932824	-0.543684	-0.370164	-0.609409
horsepower	-0.771437	0.838939	0.893646	1.000000	0.860574	-0.684259	-0.411651	-0.453669
weight	-0.831741	0.896017	0.932824	0.860574	1.000000	-0.417457	-0.306564	-0.581024
acceleration	0.420289	-0.505419	-0.543684	-0.684259	-0.417457	1.000000	0.288137	0.205873
model year	0.579267	-0.348746	-0.370164	-0.411651	-0.306564	0.288137	1.000000	0.180662
origin	0.563450	-0.562543	-0.609409	-0.453669	-0.581024	0.205873	0.180662	1.000000


Data Visualizations

Heatmap : which represents correlation between attributes

In [13]:

```
sns.heatmap(dataset.corr(),annot=True,linecolor='black', linewidths = 1)#Heatmap is a wa  
y to show some sort of matrix plot,annot is used for correlation.  
fig=plt.gcf()  
fig.set_size_inches(8,8)
```

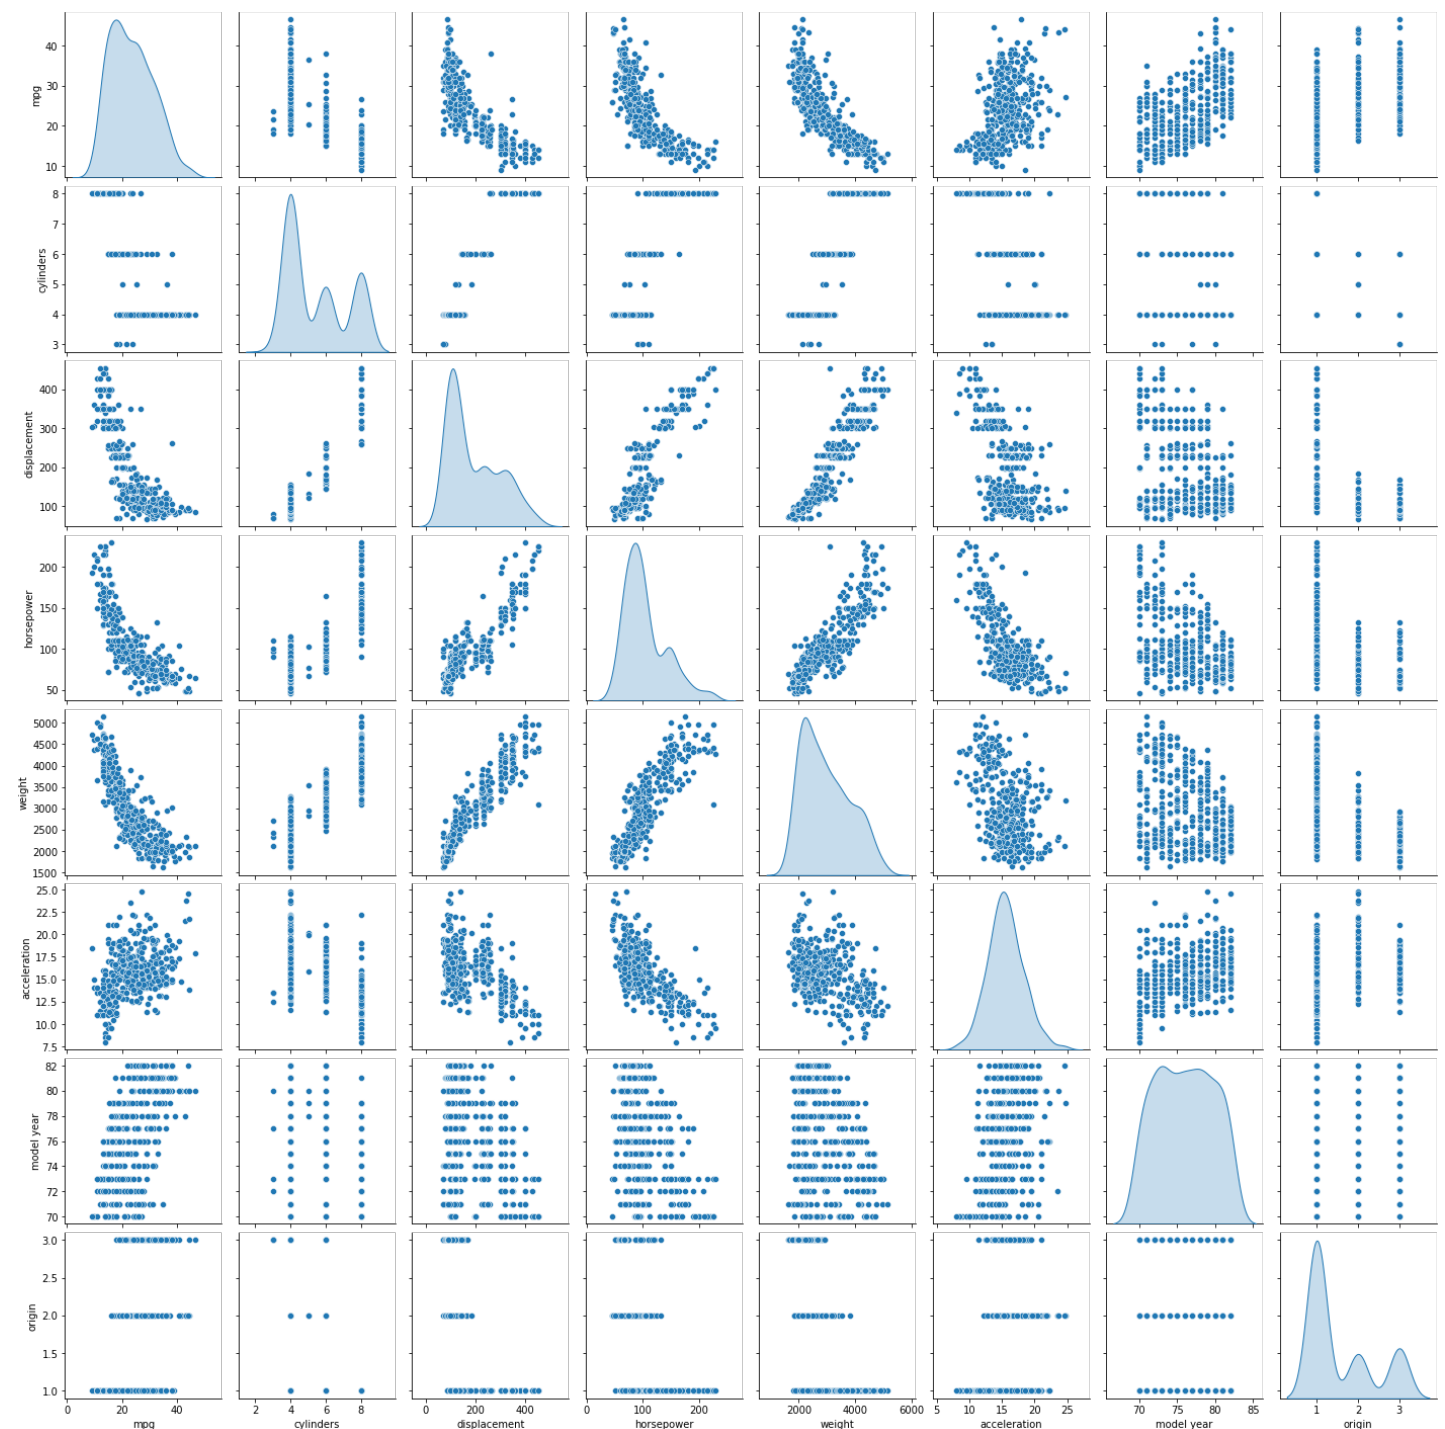




Visualizations of each attributes w.r.t rest of all attributes

In [14]:

```
sns.pairplot(dataset,diag_kind='kde') #pairplot represents pairwise relation across the entire dataframe.
plt.show()
```



Regression plots(regplot()) creates a regression line between 2 parameters and helps to visualize their linear


relationships.

In [15]:

```
sns.regplot(x="cylinders", y="mpg", data=dataset)
```

Out[15]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f092aab8750>




In [16]:

```
sns.regplot(x="displacement", y="mpg", data=dataset)
```

Out[16]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f092920a750>




In [17]:

```
sns.regplot(x="horsepower", y="mpg", data=dataset)
```

Out[17]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f092ab6c790>




In [18]:

```
sns.regplot(x="weight", y="mpg", data=dataset)
```

Out[18]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f092aa32710>




In [19]:

```
sns.regplot(x="acceleration", y="mpg", data=dataset)
```

Out[19]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f092abd9990>




In [20]:

```
sns.regplot(x="model_year", y="mpg", data=dataset)
```

Out[20]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f092aba7410>




In [21]:

```
sns.regplot(x="origin", y="mpg", data=dataset)
```

Out[21]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f092ab3c450>




In [22]:

```
sns.set(style="whitegrid")
sns.boxplot(x=dataset["mpg"])
```

Out[22]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f092ab63f10>



Finding quartiles for mpg

The P-value is the probability value that the correlation between these two variables is statistically significant.

Normally, we choose a significance level of 0.05, which means that we are 95% confident that the correlation between the variables is significant.

By convention, when the

- p-value is < 0.001 : we say there is strong evidence that the correlation is significant.
- the p-value is < 0.05 : there is moderate evidence that the correlation is significant.
- the p-value is < 0.1 : there is weak evidence that the correlation is significant.
- the p-value is > 0.1 : there is no evidence that the correlation is significant.

In [22]:


```
In [23]:
```

```
from scipy import stats
```

Cylinders vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'Cylinders' and 'mpg'.

```
In [24]:
```

```
pearson_coef, p_value = stats.pearsonr(dataset['cylinders'], dataset['mpg'])  
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

```
The Pearson Correlation Coefficient is -0.7753962854205542  with a P-value of P = 4.50399  
2246177055e-81
```

Conclusion:

Since the p-value is < 0.001 , the correlation between cylinders and mpg is statistically significant, and the coefficient of ~ -0.775 shows that the relationship is negative and moderately strong.

Displacement vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'Displacement' and 'mpg'.

```
In [25]:
```

```
pearson_coef, p_value = stats.pearsonr(dataset['displacement'], dataset['mpg'])  
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

```
The Pearson Correlation Coefficient is -0.8042028248058978  with a P-value of P = 1.65588  
89101930157e-91
```

Conclusion:

Since the p-value is < 0.1 , the correlation between displacement and mpg is statistically significant, and the linear negative relationship is quite strong (~ -0.809 , close to -1)

Horsepower vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'horsepower' and 'mpg'.

```
In [26]:
```

```
pearson_coef, p_value = stats.pearsonr(dataset['horsepower'], dataset['mpg'])  
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

```
The Pearson Correlation Coefficient is -0.7714371350025526  with a P-value of P = 9.25547  
7533166725e-80
```

Conclusion:

Since the p-value is < 0.001 , the correlation between horsepower and mpg is statistically significant, and the coefficient of ~ -0.771 shows that the relationship is negative and moderately strong.

Weight vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'weight' and 'mpg'.

```
In [28]:
```

```
pearson_coef, p_value = stats.pearsonr(dataset['weight'], dataset['mpg'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

The Pearson Correlation Coefficient is -0.831740933244335 with a P-value of P = 2.9727995640500577e-103

Conclusion:

Since the p-value is < 0.001 , the correlation between weight and mpg is statistically significant, and the linear negative relationship is quite strong (~-0.831, close to -1)

Acceleration vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'Acceleration' and 'mpg'.

In [29]:

```
pearson_coef, p_value = stats.pearsonr(dataset['acceleration'], dataset['mpg'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

The Pearson Correlation Coefficient is 0.4202889121016507 with a P-value of P = 1.823091535078553e-18

Conclusion:

Since the p-value is > 0.1 , the correlation between acceleration and mpg is statistically significant, but the linear relationship is weak (~0.420).

Model year vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'Model year' and 'mpg'.

In [30]:

```
pearson_coef, p_value = stats.pearsonr(dataset['model year'], dataset['mpg'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

The Pearson Correlation Coefficient is 0.5792671330833096 with a P-value of P = 4.844935813365483e-37

Conclusion:

Since the p-value is < 0.001 , the correlation between model year and mpg is statistically significant, but the linear relationship is only moderate (~0.579).

Origin vs mpg

Let's calculate the Pearson Correlation Coefficient and P-value of 'Origin' and 'mpg'.

In [31]:

```
pearson_coef, p_value = stats.pearsonr(dataset['origin'], dataset['mpg'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value of P =", p_value)
```

The Pearson Correlation Coefficient is 0.5634503597738431 with a P-value of P = 1.0114822102336483e-34

Conclusion:

Since the p-value is < 0.001 , the correlation between origin and mpg is statistically significant, but the linear relationship is only moderate (~0.563).

Ordinary Least Squares Statistics

In [32]:

```
test=smf.ols('mpg~cylinders+displacement+horsepower+weight+acceleration+origin',dataset).
fit()
test.summary()
```

Out[32]:

OLS Regression Results

Dep. Variable:	mpg	R-squared:	0.717
Model:	OLS	Adj. R-squared:	0.713
Method:	Least Squares	F-statistic:	165.5
Date:	Mon, 07 Nov 2022	Prob (F-statistic):	4.84e-104
Time:	09:56:04	Log-Likelihood:	-1131.1
No. Observations:	398	AIC:	2276.
Df Residuals:	391	BIC:	2304.
Df Model:	6		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	42.7111	2.693	15.861	0.000	37.417	48.005
cylinders	-0.5256	0.404	-1.302	0.194	-1.320	0.268
displacement	0.0106	0.009	1.133	0.258	-0.008	0.029
horsepower	-0.0529	0.016	-3.277	0.001	-0.085	-0.021
weight	-0.0051	0.001	-6.441	0.000	-0.007	-0.004
acceleration	0.0043	0.120	0.036	0.972	-0.232	0.241
origin	1.4269	0.345	4.136	0.000	0.749	2.105

Omnibus:	32.659	Durbin-Watson:	0.886
Prob(Omnibus):	0.000	Jarque-Bera (JB):	43.338
Skew:	0.624	Prob(JB):	3.88e-10
Kurtosis:	4.028	Cond. No.	3.99e+04

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.99e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Inference as in the above summary the p value of the acceleration is maximum(i.e 0.972) so we can remove the acc variable from the dataset

Seperating into Dependent and Independent variables

Independent variables

In [33]:

```
x=dataset[['cylinders','displacement','horsepower','weight','model year','origin']].values
x
```

Out[33]:

```
array([[8.000e+00, 3.070e+02, 1.300e+02, 3.504e+03, 7.000e+01, 1.000e+00],
       [8.000e+00, 3.500e+02, 1.650e+02, 3.693e+03, 7.000e+01, 1.000e+00],
       [8.000e+00, 3.180e+02, 1.500e+02, 3.436e+03, 7.000e+01, 1.000e+00],
       ...,
       [4.000e+00, 1.350e+02, 8.400e+01, 2.295e+03, 8.200e+01, 1.000e+00],
       [4.000e+00, 1.200e+02, 7.900e+01, 2.625e+03, 8.200e+01, 1.000e+00],
       [4.000e+00, 1.190e+02, 8.200e+01, 2.720e+03, 8.200e+01, 1.000e+00]])
```

Dependent variables

In [34]:

```
y=dataset.iloc[:,0:1].values
y
```

Out[34]:

```
array([[18. ],
       [15. ],
       [18. ],
       [16. ],
       [17. ],
       [15. ],
       [14. ],
       [14. ],
       [14. ],
       [14. ],
       [15. ],
       [15. ],
       [14. ],
       [15. ],
       [14. ],
       [24. ],
       [22. ],
       [18. ],
       [21. ],
       [27. ],
       [26. ],
       [25. ],
       [24. ],
       [25. ],
       [26. ],
       [21. ],
       [10. ],
       [10. ],
       [11. ],
       [ 9. ],
       [27. ],
       [28. ],
       [25. ],
       [25. ],
       [19. ],
       [16. ],
       [17. ],
       [19. ],
       [18. ],
       [14. ],
       [14. ],
       [14. ],
       [14. ],
       [12. ],
       [13. ],
       [13. ],
       [18. ],
       [22. ]])
```

[22.],
[19.],
[18.],
[23.],
[28.],
[30.],
[30.],
[31.],
[35.],
[27.],
[26.],
[24.],
[25.],
[23.],
[20.],
[21.],
[13.],
[14.],
[15.],
[14.],
[17.],
[11.],
[13.],
[12.],
[13.],
[19.],
[15.],
[13.],
[13.],
[14.],
[18.],
[22.],
[21.],
[26.],
[22.],
[28.],
[23.],
[28.],
[27.],
[13.],
[14.],
[13.],
[14.],
[15.],
[12.],
[13.],
[13.],
[14.],
[13.],
[12.],
[13.],
[18.],
[16.],
[18.],
[18.],
[23.],
[26.],
[11.],
[12.],
[13.],
[12.],
[18.],
[20.],
[21.],
[22.],
[18.],
[19.],
[21.],
[26.],
[15.],
[16.],
[29.],
[24.]

[21.],
[20.],
[19.],
[15.],
[24.],
[20.],
[11.],
[20.],
[21.],
[19.],
[15.],
[31.],
[26.],
[32.],
[25.],
[16.],
[16.],
[18.],
[16.],
[13.],
[14.],
[14.],
[14.],
[29.],
[26.],
[26.],
[31.],
[32.],
[28.],
[24.],
[26.],
[24.],
[26.],
[31.],
[19.],
[18.],
[15.],
[15.],
[16.],
[15.],
[16.],
[14.],
[17.],
[16.],
[15.],
[18.],
[21.],
[20.],
[13.],
[29.],
[23.],
[20.],
[23.],
[24.],
[25.],
[24.],
[18.],
[29.],
[19.],
[23.],
[23.],
[22.],
[25.],
[33.],
[28.],
[25.],
[25.],
[26.],
[27.],
[17.5],
[16.],
[15.5],
[14 5]

[11.5],
[22.],
[22.],
[24.],
[22.5],
[29.],
[24.5],
[29.],
[33.],
[20.],
[18.],
[18.5],
[17.5],
[29.5],
[32.],
[28.],
[26.5],
[20.],
[13.],
[19.],
[19.],
[16.5],
[16.5],
[13.],
[13.],
[13.],
[31.5],
[30.],
[36.],
[25.5],
[33.5],
[17.5],
[17.],
[15.5],
[15.],
[17.5],
[20.5],
[19.],
[18.5],
[16.],
[15.5],
[15.5],
[16.],
[29.],
[24.5],
[26.],
[25.5],
[30.5],
[33.5],
[30.],
[30.5],
[22.],
[21.5],
[21.5],
[43.1],
[36.1],
[32.8],
[39.4],
[36.1],
[19.9],
[19.4],
[20.2],
[19.2],
[20.5],
[20.2],
[25.1],
[20.5],
[19.4],
[20.6],
[20.8],
[18.6],
[18.1],
[19 2]

[19.2],
[17.7],
[18.1],
[17.5],
[30.],
[27.5],
[27.2],
[30.9],
[21.1],
[23.2],
[23.8],
[23.9],
[20.3],
[17.],
[21.6],
[16.2],
[31.5],
[29.5],
[21.5],
[19.8],
[22.3],
[20.2],
[20.6],
[17.],
[17.6],
[16.5],
[18.2],
[16.9],
[15.5],
[19.2],
[18.5],
[31.9],
[34.1],
[35.7],
[27.4],
[25.4],
[23.],
[27.2],
[23.9],
[34.2],
[34.5],
[31.8],
[37.3],
[28.4],
[28.8],
[26.8],
[33.5],
[41.5],
[38.1],
[32.1],
[37.2],
[28.],
[26.4],
[24.3],
[19.1],
[34.3],
[29.8],
[31.3],
[37.],
[32.2],
[46.6],
[27.9],
[40.8],
[44.3],
[43.4],
[36.4],
[30.],
[44.6],
[40.9],
[33.8],
[29.8],
[32.7],
[22 71


```

[29.7],
[35. ],
[23.6],
[32.4],
[27.2],
[26.6],
[25.8],
[23.5],
[30. ],
[39.1],
[39. ],
[35.1],
[32.3],
[37. ],
[37.7],
[34.1],
[34.7],
[34.4],
[29.9],
[33. ],
[34.5],
[33.7],
[32.4],
[32.9],
[31.6],
[28.1],
[30.7],
[25.4],
[24.2],
[22.4],
[26.6],
[20.2],
[17.6],
[28. ],
[27. ],
[34. ],
[31. ],
[29. ],
[27. ],
[24. ],
[23. ],
[36. ],
[37. ],
[31. ],
[38. ],
[36. ],
[36. ],
[36. ],
[34. ],
[38. ],
[32. ],
[38. ],
[25. ],
[38. ],
[26. ],
[22. ],
[32. ],
[36. ],
[27. ],
[27. ],
[44. ],
[32. ],
[28. ],
[31. ]]

```

Splitting into train and test data.

In [35]:

```
from sklearn.model_selection import train_test_split
```

In [36]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.1,random_state=0)
```

we are splitting as 90% train data and 10% test data

Normalisation

In []:

```
from sklearn.preprocessing import StandardScaler
sd = StandardScaler()
x_train = sd.fit_transform(x_train)
x_test = sd.fit_transform(x_test)
y_train = sd.fit_transform(y_train)
y_test = sd.fit_transform(y_test)

x_train
```

decision tree regressor

In [37]:

```
from sklearn.tree import DecisionTreeRegressor
dt=DecisionTreeRegressor(random_state=0,criterion="mae")
dt.fit(x_train,y_train)
```

/usr/local/lib/python3.7/dist-packages/sklearn/tree/_classes.py:370: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in version 1.2. Use `criterion='absolute_error'` which is equivalent.

FutureWarning,

Out[37]:

DecisionTreeRegressor(criterion='mae', random_state=0)

In [38]:

```
import pickle
pickle.dump(dt,open('decision_model.pkl','wb'))
```

In [39]:

```
y_pred=dt.predict(x_test)
y_pred
```

Out[39]:

```
array([14. , 26.5, 14. , 19. , 18. , 31. , 37. , 22. , 15. , 26.8, 34.5,
       31.8, 16. , 26. , 15.5, 34.3, 26.6, 27. , 16. , 31.5, 16. , 23. ,
       27.5, 19. , 33.8, 24.2, 36.1, 36. , 34.3, 18.5, 19.2, 35. , 15. ,
       32. , 22. , 23. , 19.4, 16. , 31.5, 12. ])
```

In [40]:

```
y_test
```

Out[40]:

```
array([[14. ],
       [25. ],
       [13. ],
       [21. ],
       [18. ],
       [35. ],
       [34.1],
       [20. ],
       [15. ],
```

```
[23.5],
[40.9],
[37.2],
[18. ],
[23. ],
[15.5],
[35.7],
[31. ],
[27. ],
[18. ],
[37.3],
[15.5],
[23. ],
[24. ],
[18. ],
[34.5],
[25.4],
[36.1],
[34. ],
[30. ],
[16. ],
[18.6],
[37. ],
[15. ],
[33.5],
[22.4],
[24. ],
[19. ],
[16.9],
[31.9],
[12. ]])
```

In [60]:

```
ax1 = sns.distplot(dataset['mpg'], hist=False, color="r", label="Actual Value")
sns.distplot(y_pred, hist=False, color="b", label="Fitted Values" , ax=ax1)
```

```
plt.title('Actual vs Fitted Values for mpg')
plt.xlabel('mpg')
plt.ylabel('Proportion of Cars')
```


```
plt.show()
plt.close()
```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

warnings.warn(msg, FutureWarning)

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

warnings.warn(msg, FutureWarning)



We can see that the fitted values are reasonably close to the actual values, since the two distributions overlap a bit. However, there is definitely some room for improvement.

R-squared

R-squared is a statistical measure of how close the data are to the fitted regression line. It is also known as the coefficient of determination, or the coefficient of multiple determination for multiple regression.

R-squared = Explained variation / Total variation

Mean Squared Error (MSE)

The Mean Squared Error measures the average of the squares of errors, that is, the difference between actual value (y) and the estimated value (\hat{y}).

In [42]:

```
from sklearn.metrics import r2_score, mean_squared_error
```

In [43]:

```
r2_score(y_test, y_pred)
```

Out[43]:

```
0.912578781275149
```

In [44]:

```
mean_squared_error(y_test, y_pred)
```

Out[44]:

```
6.042499999999999
```

In [45]:

```
np.sqrt(mean_squared_error(y_test, y_pred))
```

Out[45]:

```
2.458149710656371
```

random forest regressor

In [46]:

```
from sklearn.ensemble import RandomForestRegressor
```

In [47]:

```
rf = RandomForestRegressor(n_estimators=10, random_state=0, criterion='mae')  
rf.fit(x_train, y_train)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/ensemble/_forest.py:407: FutureWarning: Cr  
iterion 'mae' was deprecated in v1.0 and will be removed in version 1.2. Use `criterion='  
absolute_error` which is equivalent.
```

```
FutureWarning,
```

Out[47]:

```
RandomForestRegressor(criterion='mae', n_estimators=10, random_state=0)
```

In [48]:

```
y_pred2=rf.predict(x_test)
y_pred2
```

Out[48]:

```
array([14.05, 25.55, 13.7 , 21.2 , 18.5 , 30.8 , 34.31, 22.5 , 15.1 ,
       24.46, 32.01, 39.79, 17.8 , 24.85, 15.85, 31.31, 28.32, 26.93,
       16.54, 32.12, 16.05, 25.6 , 23.86, 20.56, 32.19, 24.75, 32.39,
       32.64, 31.02, 16.39, 18.35, 30.1 , 17.67, 31. , 22.91, 23.5 ,
       19.77, 16.1 , 33.65, 12.  ])
```

In [49]:

```
ax1 = sns.distplot(dataset['mpg'], hist=False, color="r", label="Actual Value")
sns.distplot(y_pred2, hist=False, color="b", label="Fitted Values" , ax=ax1)
```

```
plt.title('Actual vs Fitted Values for mpg')
plt.xlabel('mpg')
plt.ylabel('Proportion of Cars')
```

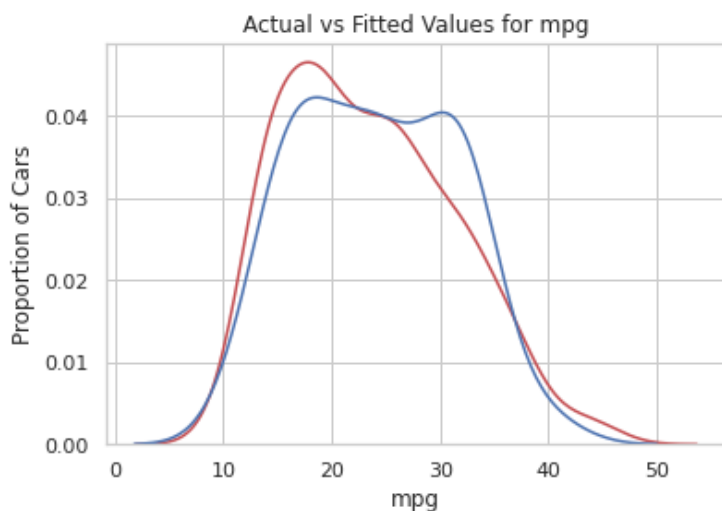
```
plt.show()
plt.close()
```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

warnings.warn(msg, FutureWarning)

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

warnings.warn(msg, FutureWarning)



We can see that the fitted values are reasonably close to the actual values, since the two distributions overlap a bit. However, there is definitely some room for improvement.

In [50]:

```
from sklearn.metrics import r2_score, mean_squared_error
```

In [51]:

```
r2_score(y_test, y_pred2)
```

Out[51]:

```
0.9013457876319049
```

In [52]:

```
mean_squared_error(y_test, y_pred2)
```

```
mean_squared_error(y_test,y_pred2)
```

Out[52]:

```
6.8189175000000002
```

In [53]:

```
np.sqrt(mean_squared_error(y_test,y_pred2))
```

Out[53]:

```
2.611305707878724
```

linear regression

In [54]:

```
from sklearn.linear_model import LinearRegression
mr=LinearRegression()
mr.fit(x_train,y_train)
```

Out[54]:

```
LinearRegression()
```

In [55]:

```
y_pred3=mr.predict(x_test)
y_pred3
```

Out[55]:

```
array([[13.20818031],
       [24.27993342],
       [11.61339788],
       [20.96914745],
       [17.7247275 ],
       [29.44595217],
       [33.47372984],
       [23.1855594 ],
       [15.045202  ],
       [26.79998444],
       [32.32754229],
       [33.93400668],
       [21.48572281],
       [25.80404696],
       [16.32002867],
       [30.62069212],
       [28.3611479 ],
       [28.68598061],
       [17.66367225],
       [31.02921296],
       [15.54781059],
       [24.61489613],
       [26.90655487],
       [20.51716586],
       [29.66216351],
       [28.48379869],
       [31.00137585],
       [29.9752557 ],
       [29.90123742],
       [18.07465439],
       [20.36226872],
       [31.32907003],
       [20.95979818],
       [32.03796407],
       [23.8731354 ],
       [26.30724058],
       [21.37158555],
       [16.80870416],
       [32.14991802],
```

```
[ 9.27600756]])
```

In [56]:

```
ax1 = sns.distplot(dataset['mpg'], hist=False, color="r", label="Actual Value")
sns.distplot(y_pred3, hist=False, color="b", label="Fitted Values" , ax=ax1)

plt.title('Actual vs Fitted Values for mpg')
plt.xlabel('mpg')
plt.ylabel('Proportion of Cars')


plt.show()
plt.close()
```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

warnings.warn(msg, FutureWarning)

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

warnings.warn(msg, FutureWarning)



We can see that the fitted values are not as close to the actual values, since the two distributions overlap a bit. However, there is definitely some room for improvement.

In [57]:

```
from sklearn.metrics import r2_score, mean_squared_error
r2_score(y_test, y_pred3)
```

Out[57]:

0.8460443802521529

In [58]:

```
mean_squared_error(y_test, y_pred3)
```

Out[58]:

10.64131621470885

In [59]:

```
np.sqrt(mean_squared_error(y_test, y_pred3))
```

Out[59]:

3.262103035575187

Conclusion:

When comparing models, the model with the higher R-squared value is a better fit for the data.

When comparing models, the model with the smallest MSE value is a better fit for the data.

Comparing these three models, we conclude that the DecisionTree model is the best model to be able to predict mpg from our dataset.